




# Advanced SQL II



THE KNOWLEDGE HOUSE

## Agenda - Schedule

1. SQL Leetcode Q
2. Understanding Entity Relationship Diagrams
3. Using SQLite in Python
4. Break
5. SQLite Lab



*Database systems of the past*



## Agenda - Announcements

- No pre-class quiz
- TLAB #3 due 5/14
  - Early grade due date: 5/7
  - Extension due date: 5/13
- In-class end of phase project is being released **THIS THURSDAY!** (We recommend attending this review session)



## Agenda - Goals

- Interpret ERD Diagrams.
- Connect to a SQLite database using Python.
- Create a table and insert data into it.
- Run simple SQL queries from Python to explore data.
- Print and interpret results returned by the database.

# SQL Leetcode Q

---

## 584. Find Customer Referee

Easy

Topics

Companies

Hint

[SQL Schema](#) > [Pandas Schema](#) >

Table: Customer

Column Name	Type
id	int
name	varchar
referee_id	int

In SQL, id is the primary key column for this table.

Each row of this table indicates the id of a customer, their name, and the id of the customer who referred them.

Take 10 minutes to complete “Find Customer Referee”:

<https://leetcode.com/problems/find-customer-referee/description/?envType=study-plan-v2&envId=top-sql-50>

# Entity Relationship Diagrams (ERDs)

---

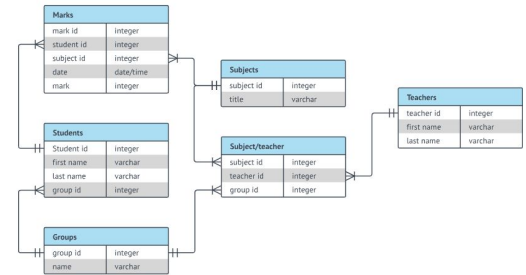
# Entity Relationship Diagram (ERD)

This previous week, we learned that we often store **multiple tables** in one database.

Like we discussed, we abide by this structure due to the principles of **normalization** to ensure that we have **data integrity**.

One way to visualize the relationship of these tables is through an *Entity Relationship Diagram*

Why might we want to keep track of the structure of our database visually? What might this help us with?



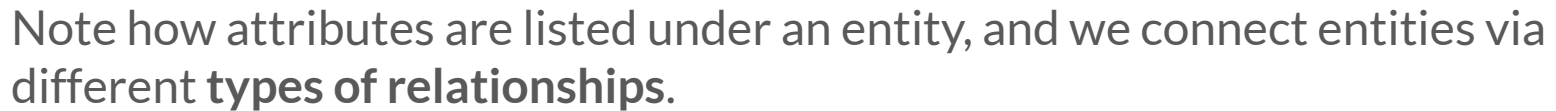




## What is an ERD?

An **Entity Relationship Diagram (ERD)** is a visual representation of the **relationships among entities in a database**. Let's iron out some definitions of an ERD:


- **Entities:** Represented by **rectangles** and are the **objects we are interested in keeping information on** (e.g., *Customer*, *Product*, *Student*). **Rows of a table.**
- **Attributes:** **Characteristics** of entities (e.g., *CustomerID*, *ProductName*). **Columns of a table.**
- **Relationships:** **Connections** between entities, indicating how they are related.



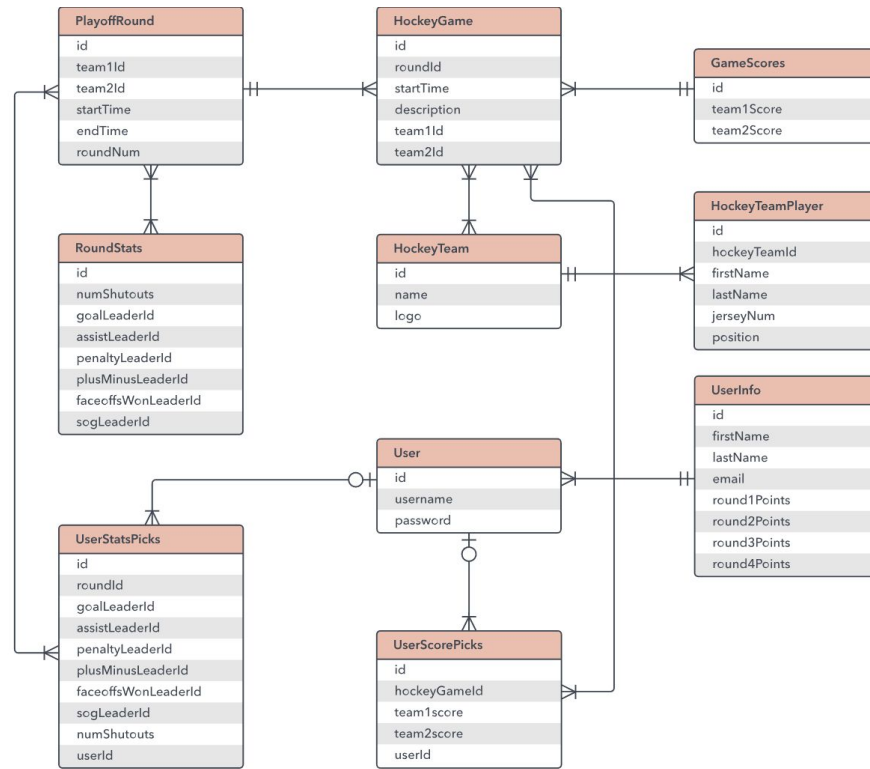
# ERD Relationships/Cardinality

While **notation** for these types of relationships vary, they generally describe the following:

- **One-to-One (Zero or one):** **One** entity is strictly associated with **one** entity (ex: *one SSN per person, and one person per SSN*)
- **One-to-Many:** **One** object can be associated with **many** other objects (ex: *one person can have multiple orders, but one order can only have one person*)
- **Many-to-Many:** **Multiple** objects can be associated with **multiple** other objects (ex: *one student can take multiple any courses, and one course can have multiple students*)



See if you could interpret the relationship between other entities, such as HockeyTeam to HockeyGame; GameScores to HockeyGame; HockeyTeam to HockeyTeamPlayer; HockeyTeam to HockeyTeamPlayer.



For example, in this diagram the **User to UserStatsPicks** entities is **one (or zero) to many**. In other words, **one user can have many stats picks**, but **one stats picks can only have one user**.

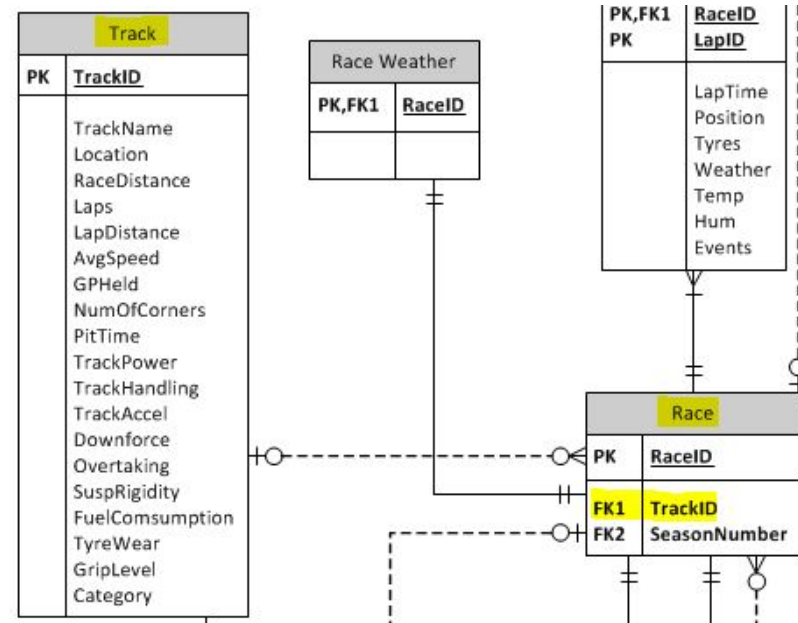
## ERD Keys

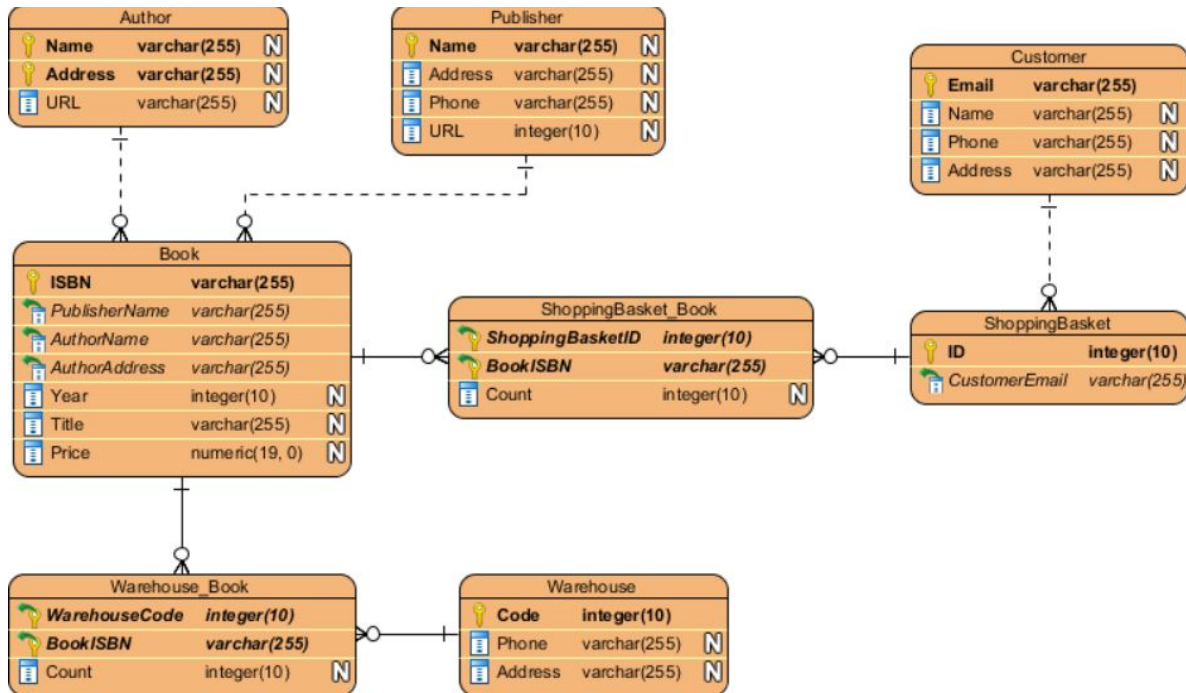
To sufficiently discuss ERD diagrams, we should also continue our discussion of **primary/foreign keys**.

**Primary Keys** are a unique identifier for each object in a table. It distinguishes one record from another within the same relation.

**Foreign Keys** are a column in a database table that is used to establish a link between two tables. It creates a relationship by referencing the primary key of another table.

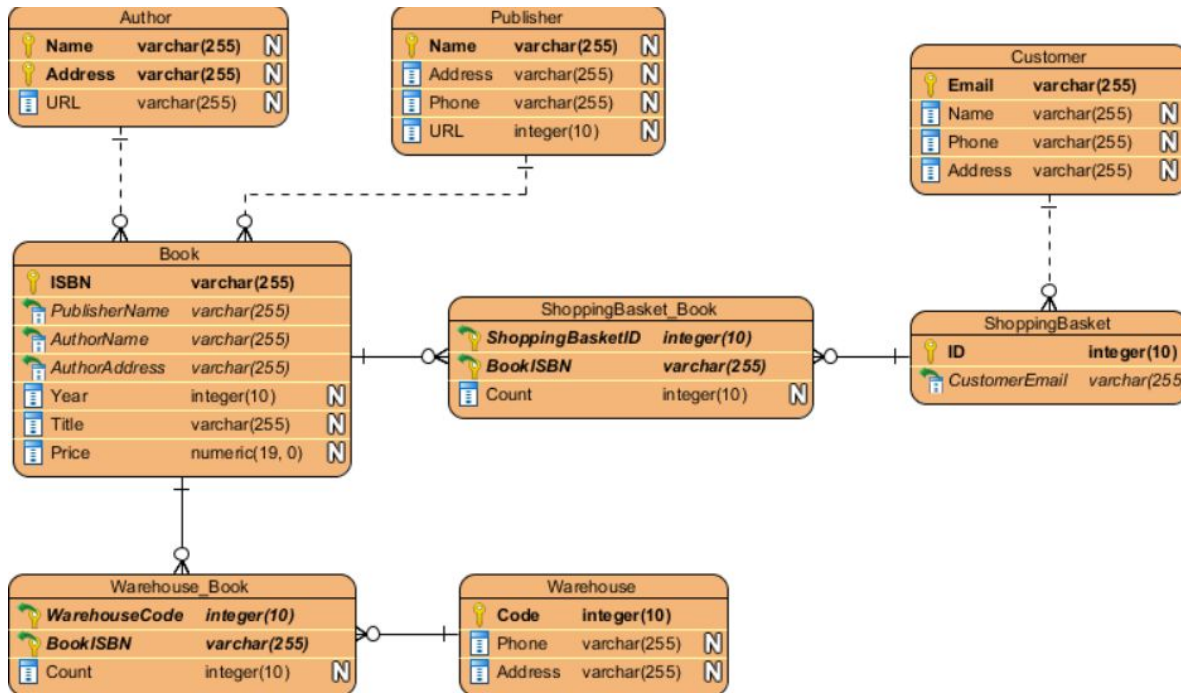
Note that TrackID shows up as a Primary Key in the “Track” table and a foreign key in the “Race” table.






ERD Diagrams allow us to document and understand the structure of a database. Additionally it also allows us to anticipate the types of JOINS that we will perform when we are querying our database.


For example, let's take a look at the ERD diagram above. If we would like to get all available information on a **Book** and its **Author**, which tables should we join and on which columns should we be performing this join?



**SELECT \***  
**FROM Book b JOIN Author a**  
**ON b.AuthorName = a.Name**  
**;**



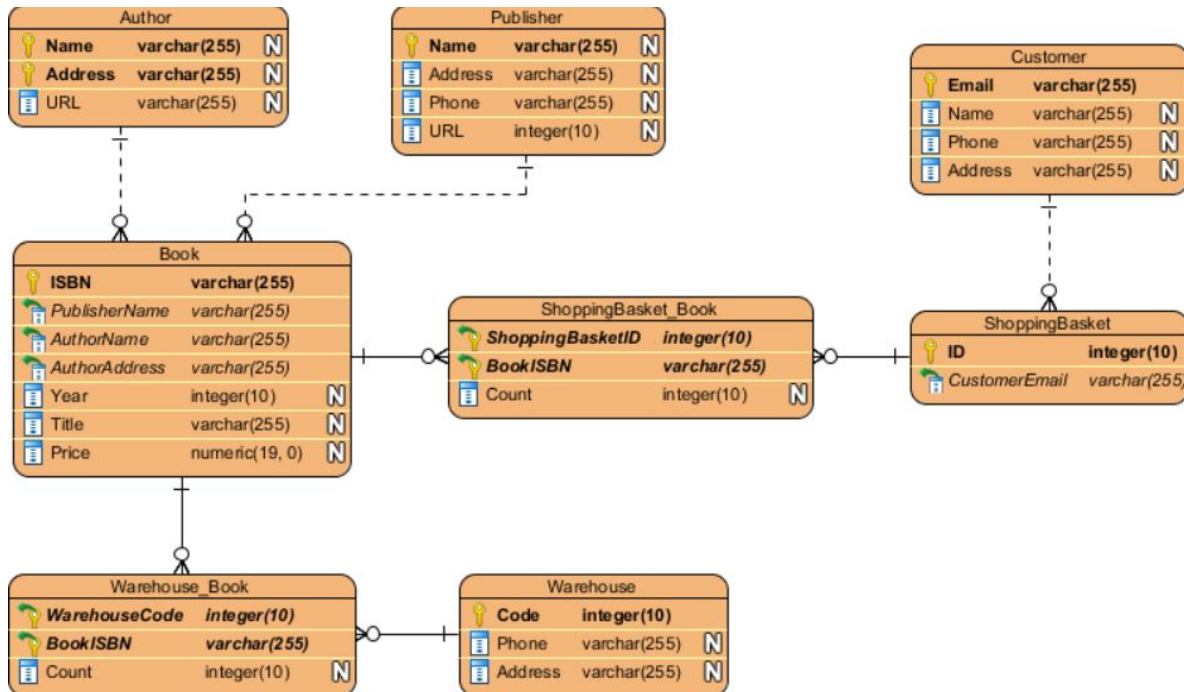
We can join the **Book** and **Author** table on the “b.AuthorName” foreign key and “a.Name” primary key. Using the ERD diagram relationship symbols, we can also state the relationship between these two tables. **Can one author have more than one book?**



```
SELECT *
FROM Book b JOIN Author a
ON b.AuthorName = a.Name
;
```

Yes! But also, note that an author can have **zero books** as well. If we want to preserve authors who have zero books in this query, **what kind of join can we use?**





```

SELECT *
FROM Book b LEFT JOIN Author a
ON b.AuthorName = a.Name
;
  
```

As we discussed in previous lessons, the **LEFT JOIN** ensures that we keep Authors who have not published any books (aka rows that result in null values).

# Using SQLite in Python

---

## SQLite + Python



Like we established last week, relational database management is tantamount for contemporary tech-stacks as they provide a **central system for storing & querying data**. This includes:

- *Storage of massive amounts of data*
- *Access via a query language*
- *Durability even during power failures*
- *Allow multiple users to interact with data*

But how do we actually go about implementing this technology in the context of everything we've learned so far? To accomplish this, we will use **SQLite**.



Like we established, there exist **different flavors** of SQL. We might opt for PostgreSQL for vector-database support, MySQL for typical LAMP applications. However SQLite is a perfect package **for non-server database based approaches.**

## SQLite + Python



While almost all other SQL approaches require you to **set up a database server**, **SQLite** is special in the sense **that is natively supported in Python and stores the database within a lightweight local file.**

Simply import the **sqlite3** package in your Python module to get started!

Before we run through an example, let's consider if a **Python module** or **Jupyter notebook** is appropriate for this exercise...



## Choosing the Right Tool


Remember! Part of being a good technologist is being able to discern which tools are right for a specific task.

Follow this rule of thumb:

- **Python Modules:** Any functionality that should be run automatically (on a schedule) or integrated into a larger system.
- **Jupyter Notebooks:** Making tutorials, reports, or code that you will share with a non-technologist stakeholder.

Given these slides from “Applied Rest APIs I”, which tool should we opt for in this case?

# ETL Process



We will often integrate SQL-adjacent packages in our ETL pipelines which should be run on a **scheduled and consistent basis**. Therefore, let's utilize **Python modules** for this exercise.

```
import sqlite3
```

```
conn = sqlite3.connect("Fellowship.db")
```

```
cursor = conn.cursor()
```

```
cursor.execute("SELECT * FROM fellows;")
```

Within the `sqlite_intro` folder let's access `query_db.py` to find out how we can run SQL in Python.



```
import sqlite3
```

```
conn = sqlite3.connect("Fellowship.db")
```

```
cursor = conn.cursor()
```

```
cursor.execute("SELECT * FROM fellows;")
```

We establish our connection to the database using the `connect()` method. We supply the path to an already created database file called "Fellowship.db."

```
import sqlite3
```

```
conn = sqlite3.connect("Fellowship.db")
```

```
cursor = conn.cursor()
```

```
cursor.execute("SELECT * FROM fellows;")
```

Even though we established a connection, we still cannot read from the database until we create a **cursor**. This object allows us to **execute SQL statements** and **fetch** results from the database itself.

```
import sqlite3
```

```
conn = sqlite3.connect("Fellowship.db")
```

```
cursor = conn.cursor()
```

```
cursor.execute("SELECT * FROM fellows;")
```

Additionally note that you must be confident that your SQL statement is correct before writing out your query. This adds an additional challenge to our development.

Finally, we can begin **sending** SQL statements in the form of a string to our database via the `execute()` method. **Note that this requires a bit of context-switching** (sort of like inserting Russian statements in the middle of an essay written in English). This is **generally considered a bad-practice**, but we will opt for this for lack of better tools (for now).

```
cursor.execute("""  
    SELECT *  
    FROM fellows;  
""")
```

```
rows = cursor.fetchall()  
print(rows)
```

We can make our SQL statements look like idiomatic SQL by introducing triple-quoted strings. This allows us to write strings across multiple lines.

```
cursor.execute("""  
    SELECT *  
    FROM fellows;  
""")
```

```
rows = cursor.fetchall()  
print(rows)
```

```
[(1, 'Alice', 'Adams', 2, 32, 14.31), (2, 'Brian', 'Adams', 2, 32, 14.31), (3, 'Alice', 'Adams', 2, 32, 14.31), (4, 'Alice', 'Adams', 2, 32, 14.31), (5, 'Alice', 'Adams', 2, 32, 14.31), (6, 'Farrah', 'Foster', 2, 22, 27.33), (7, 'Farrah', 'Foster', 2, 22, 27.33), (8, 'Farrah', 'Foster', 2, 22, 27.33), (9, 'Farrah', 'Foster', 2, 22, 27.33), (10, 'Farrah', 'Foster', 2, 22, 27.33), (11, 'Kim', 'Klein', 1, 3, 28.71), (12, 'Kim', 'Klein', 1, 3, 28.71), (13, 'Kim', 'Klein', 1, 3, 28.71), (14, 'Omar', 'Owens', 3, 32, 10.29), (15, 'Omar', 'Owens', 3, 32, 10.29), (16, 'Priya', 'Adams', 1, 32, 21.22), (17, 'Priya', 'Adams', 1, 32, 21.22), (18, 'Priya', 'Adams', 1, 32, 21.22), (19, 'Priya', 'Adams', 1, 32, 21.22), (20, 'Priya', 'Adams', 1, 32, 21.22)]
```

More importantly however, let's focus on the output that is provide from our the statement that we just executed. Note that we get the **entire output** using the **fetchall()** method. We can have the result of this method call in a variable and print it out. **What kind of data-structure does this provide?**

```
cursor.execute("""  
    SELECT *  
    FROM fellows;  
""")
```

```
rows = cursor.fetchall()  
print(rows)
```


```
[(1, 'Alice', 'Adams', 2, 32, 14.31), (2, 'Brian', 'Adams', 2, 32, 14.31), (3, 'Alice', 'Adams', 2, 32, 14.31), (4, 'Alice', 'Adams', 2, 32, 14.31), (5, 'Alice', 'Adams', 2, 32, 14.31), (6, 'Farrah', 'Foster', 2, 22, 27.33), (7, 'Farrah', 'Foster', 2, 22, 27.33), (8, 'Farrah', 'Foster', 2, 22, 27.33), (9, 'Farrah', 'Foster', 2, 22, 27.33), (10, 'Farrah', 'Foster', 2, 22, 27.33), (11, 'Kim', 'Klein', 1, 3, 10.29), (12, 'Kim', 'Klein', 1, 3, 10.29), (13, 'Kim', 'Klein', 1, 3, 10.29), (14, 'Omar', 'Owens', 3, 32, 10.29), (15, 'Omar', 'Owens', 3, 32, 10.29), (16, 'Priya', 'Adams', 1, 32, 10.29), (17, 'Priya', 'Adams', 1, 32, 10.29), (18, 'Priya', 'Adams', 1, 32, 10.29), (19, 'Priya', 'Adams', 1, 32, 10.29), (20, 'Priya', 'Adams', 1, 32, 10.29)]
```

Note that we have a **list of tuples**! Printing out the entire list is **insufficient to properly view the data**. If we wanted to **iterate** through each element of the list, and **print out each tuple on a separate line**, what kind of Python syntax can we write?

```
cursor.execute("""  
    SELECT *  
    FROM fellows;  
""")
```

```
rows = cursor.fetchall()  
for r in rows:  
    print(r)
```

```
(1, 'Alice', 'Adams', 2, 32, 14.31)  
(2, 'Brian', 'Brown', 2, 27, 22.26)  
(3, 'Carmen', 'Chen', 1, 38, 18.04)  
(4, 'David', 'Diaz', 1, 21, 5.19)  
(5, 'Ella', 'Evans', 2, 31, 28.15)  
(6, 'Farrah', 'Foster', 2, 22, 27.33)
```



Since we save our list to a variable, we could instead iterate through this list using a for loop to cleanly print out our data!

```
cursor.execute("""  
    ...  
""")
```

```
rows = cursor.fetchall()  
print(rows)
```

Now that we understand the syntax of running SQL statements in Python, let's try a few exercises out! **How can we output all fellows who spent less than 10 hours on Canvas using a SQL statement?** Look to the triple-quoted string in the Python module above to find out more about the columns.



# Making Changes to a DB in Python

---



## Making Changes

While we went over running simple DML queries using `sqlite3`, how can we effectively **commit** a change to our database from Python?

Before we go over this syntax, let's recall the types of DDL queries that we can make, starting with the **CREATE** clause.

Statement	Description
CREATE	Create a new object in the database, such as a table or a view.
ALTER	Modify the structure of an object. For instance, altering a table to add a new column.
DROP	Remove an object from the database.
RENAME	Rename an existing object.

Taken from the [Azure DP-900 prep](#)



## SQL - Logical Groups - DDL

```
CREATE TABLE Product (  
    ID INT PRIMARY KEY,  
    NAME VARCHAR(20) NOT NULL,  
    PRICE DECIMAL DEFAULT 0.0  
);
```

**PRIMARY\_KEY:** uniquely identifies the sample (row)

**FOREIGN\_KEY:** this row is connected to another table

**NOT NULL** means the value must NOT be empty (null)

**DEFAULT** gives a default value if nothing is put in during insertion

```
cursor.execute("""  
    CREATE TABLE IF NOT EXISTS staff (  
        staffid INTEGER PRIMARY KEY,  
        first_name TEXT,  
        last_name TEXT,  
        trackid INTEGER,  
        age INTEGER  
    );  
""")
```

```
conn.commit()
```

Given the previous DDL statement, we can simply include the CREATE TABLE clause in a triple-quoted string within the **execute** method, but note that we did not actually make any changes to our database just yet!

```
cursor.execute("""  
    CREATE TABLE IF NOT EXISTS staff (  
        staffid INTEGER PRIMARY KEY,  
        first_name TEXT,  
        last_name TEXT,  
        trackid INTEGER,  
        age INTEGER  
    );  
""")
```

```
conn.commit()
```

Just like with git, we must **commit** any DDL query to the connection if we would like these changes to be reflected! Without a commit the database is left **unmodified**.



# SQLite Lab

---



## SQL

Complete the module `lab_part1.py` and `lab_part2.py` modules for the remainder of the lab time





# Wednesday

## SQL Leetcode Review

- *Another Leetcode day! (pray that it doesn't go down)*
- *We'll hold off on Machine Learning until Phase 2*

